



IBM Developer
SKILLS NETWORK

(https://cognitiveclass.ai/?utm_medium=Exinfluencer&utm_source=Exinfluencer&utm_content=000026UJ&utm_term=10006555&utm_id=NASkillsNetwork-Channel-SkillsNetworkCoursesIBMDeveloperSkillsNetworkRP0321ENSkillsNetwork25371262-2021-01-01)

Predict Hourly Rented Bike Count using Basic Linear Regression Models

Estimated time needed: **90** minutes

Lab Overview:

Now that you have performed exploratory analysis on the bike sharing demand dataset and obtained some insights on the attributes, it's time to build predictive models to predict the hourly rented bike count using related weather and date information.

In this lab, you will be asked to use `tidymodels` to build some baseline linear regression models:

- **TASK: Split data into training and testing datasets**
- **TASK: Build a linear regression model using only the weather variables**
- **TASK: Build a linear regression model using both weather and date variables**
- **TASK: Evaluate the models and identify important variables**

Let's start!

First install and import the necessary libraries

```
In [3]: # It may take several minutes to install those libraries in Watson Studio
install.packages("rlang")
install.packages("tidymodels")
```

```
Updating HTML index of packages in '.Library'
Making 'packages.html' ... done
also installing the dependencies 'listenv', 'parallelly', 'future', 'warp',
'lhs', 'DiceDesign', 'glue', 'lifecycle', 'tidyselect', 'vctrs', 'pillar', 'i
soband', 'patchwork', 'generics', 'globals', 'ipred', 'frrrr', 'slider', 'ell
ipsis', 'cpp11', 'GPfit', 'pROC', 'broom', 'cli', 'conflicted', 'dials', 'dpl
yr', 'ggplot2', 'hardhat', 'infer', 'modeldata', 'parsnip', 'purrr', 'recipe
s', 'rsample', 'rstudioapi', 'tibble', 'tidyr', 'tune', 'workflows', 'workflo
wsets', 'yardstick'
```

```
Updating HTML index of packages in '.Library'
Making 'packages.html' ... done
```

```
In [4]: library("tidymodels")
library("tidyverse")
library("stringr")
```

```
Registered S3 method overwritten by 'tune':
  method          from
  required_pkgs.model_spec parsnip
— Attaching packages ————— tidymodels 0.1.4
—
```

✓ broom	0.7.9	✓ recipes	0.1.17
✓ dials	0.0.10	✓ rsample	0.1.0
✓ dplyr	1.0.7	✓ tibble	3.1.5
✓ ggplot2	3.3.5	✓ tidyr	1.1.4
✓ infer	1.0.0	✓ tune	0.1.6
✓ modeldata	0.1.1	✓ workflows	0.2.4
✓ parsnip	0.1.7	✓ workflowsets	0.1.0
✓ purrr	0.3.4	✓ yardstick	0.0.8

```
— Conflicts ————— tidymodels_conflicts()
—
```

```
✗ purrr::discard() masks scales::discard()
✗ dplyr::filter() masks stats::filter()
✗ dplyr::lag() masks stats::lag()
✗ recipes::step() masks stats::step()
• Dig deeper into tidy modeling with R at https://www.tmr.org
```

```
Registered S3 method overwritten by 'rvest':
  method          from
  read_xml.response xml2
— Attaching packages ————— tidyverse 1.2.1
—
```

✓ readr	1.3.1	✓ stringr	1.4.0
✓ readr	1.3.1	✓ forcats	0.4.0

```
— Conflicts ————— tidyverse_conflicts()
—
```

```
✗ readr::col_factor() masks scales::col_factor()
✗ purrr::discard() masks scales::discard()
✗ dplyr::filter() masks stats::filter()
✗ stringr::fixed() masks recipes::fixed()
✗ dplyr::lag() masks stats::lag()
✗ readr::spec() masks yardstick::spec()
```

The `seoul_bike_sharing_converted_normalized.csv` will be our main dataset which has following variables:

The response variable:

- RENTED_BIKE_COUNT - Count of bikes rented at each hour

Weather predictor variables:

- TEMPERATURE - Temperature in Celsius
- HUMIDITY - Unit is %
- WIND_SPEED - Unit is m/s
- VISIBILITY - Multiplied by 10m
- DEW_POINT_TEMPERATURE - The temperature to which the air would have to cool down in order to reach saturation, unit is Celsius
- SOLAR_RADIATION - MJ/m²
- RAINFALL - mm
- SNOWFALL - cm

Date/time predictor variables:

- DATE - Year-month-day
- HOUR - Hour of the day
- FUNCTIONAL_DAY - NoFunc(Non Functional Hours), Fun(Functional hours)
- HOLIDAY - Holiday/No holiday
- SEASONS - Winter, Spring, Summer, Autumn

Let's read the dataset as a dataframe first:

```
In [5]: # Dataset URL
dataset_url <- "https://cf-courses-data.s3.us.cloud-object-storage.appdomain.c
loud/IBMDeveloperSkillsNetwork-RP0321EN-SkillsNetwork/labs/datasets/seoul_bike
_sharing_converted_normalized.csv"
bike_sharing_df <- read_csv(dataset_url)
spec(bike_sharing_df)
```

Parsed with column specification:

```
cols(  
  .default = col_double(),  
  DATE = col_character(),  
  FUNCTIONING_DAY = col_character()  
)
```

See spec(...) for full column specifications.

```
cols(  
  DATE = col_character(),  
  RENTED_BIKE_COUNT = col_double(),  
  TEMPERATURE = col_double(),  
  HUMIDITY = col_double(),  
  WIND_SPEED = col_double(),  
  VISIBILITY = col_double(),  
  DEW_POINT_TEMPERATURE = col_double(),  
  SOLAR_RADIATION = col_double(),  
  RAINFALL = col_double(),  
  SNOWFALL = col_double(),  
  FUNCTIONING_DAY = col_character(),  
  `0` = col_double(),  
  `1` = col_double(),  
  `10` = col_double(),  
  `11` = col_double(),  
  `12` = col_double(),  
  `13` = col_double(),  
  `14` = col_double(),  
  `15` = col_double(),  
  `16` = col_double(),  
  `17` = col_double(),  
  `18` = col_double(),  
  `19` = col_double(),  
  `2` = col_double(),  
  `20` = col_double(),  
  `21` = col_double(),  
  `22` = col_double(),  
  `23` = col_double(),  
  `3` = col_double(),  
  `4` = col_double(),  
  `5` = col_double(),  
  `6` = col_double(),  
  `7` = col_double(),  
  `8` = col_double(),  
  `9` = col_double(),  
  AUTUMN = col_double(),  
  SPRING = col_double(),  
  SUMMER = col_double(),  
  WINTER = col_double(),  
  HOLIDAY = col_double(),  
  NO_HOLIDAY = col_double()  
)
```

We won't be using the `DATE` column, because 'as is', it basically acts like an data entry index. (However, given more time, we could use the `DATE` column to create a 'day of week' or 'isWeekend' column, which we might expect has an affect on preferred bike rental times.) We also do not need the `FUNCTIONAL DAY` column because it only has one distinct value remaining (`YES`) after missing value processing.

```
In [6]: bike_sharing_df <- bike_sharing_df %>%  
        select(-DATE, -FUNCTIONING_DAY)
```

TASK: Split training and testing data

First, we need to split the full dataset into training and testing datasets.

The training dataset will be used for fitting regression models, and the testing dataset will be used to evaluate the trained models.

TODO: Use the `initial_split()` , `training()` , and `testing()` functions to generate a training dataset consisting of 75% of the original dataset, and a testing dataset using the remaining 25%.

```
In [7]: # Use the `initial_split()`, `training()`, and `testing()` functions to split
        # the dataset
        # With seed 1234
        set.seed(1234)

        # prop = 3/4
        # train_data
        # test_data

        data_split <- initial_split(bike_sharing_df, prop = .75)

        train_data <- training(data_split)
        test_data  <- testing(data_split)

        head(train_data)

        head(test_data)
```

RENTED_BIKE_COUNT	TEMPERATURE	HUMIDITY	WIND_SPEED	VISIBILITY	DEW_POINT_TEMPERATURE
2716	0.6066434	0.4795918	0.18918919	0.8160162	0.5195134
419	0.3776224	0.7040816	0.00000000	0.9974658	0.5195134
1106	0.5716783	0.5102041	0.45945946	0.9944247	0.5195134
154	0.4178322	0.7346939	0.06756757	0.5195134	0.5195134
192	0.5000000	0.6020408	0.06756757	1.0000000	0.5195134
213	0.1853147	0.3877551	0.41891892	1.0000000	0.5195134

RENTED_BIKE_COUNT	TEMPERATURE	HUMIDITY	WIND_SPEED	VISIBILITY	DEW_POINT_TEMPERATURE
100	0.1993007	0.3775510	0.2027027	1.0000000	0.5195134
181	0.1958042	0.3571429	0.1756757	1.0000000	0.5195134
360	0.3024476	0.2142857	0.1756757	0.9675621	0.5195134
555	0.3251748	0.5918367	0.2162162	1.0000000	0.5195134
600	0.3111888	0.7857143	0.2297297	1.0000000	0.5195134
405	0.2972028	0.8265306	0.1081081	0.8413583	0.5195134

TASK: Build a linear regression model using weather variables only

As you could imagine, weather conditions may affect people's bike renting decisions. For example, on a cold and rainy day, you may choose alternate transportation such as a bus or taxi. While on a nice sunny day, you may want to rent a bike for a short-distance travel.

Thus, can we predict a city's bike-sharing demand based on its local weather information? Let's try to build a regression model to do that.

TODO: Build a linear regression model called `lm_model_weather` using the following variables:

- TEMPERATURE - Temperature in Celsius
- HUMIDITY - Unit is %
- WIND_SPEED - Unit is m/s
- VISIBILITY - Multiplied by 10m
- DEW_POINT_TEMPERATURE - The temperature to which the air would have to cool down in order to reach saturation, unit is Celsius
- SOLAR_RADIATION - MJ/m2
- RAINFALL - mm
- SNOWFALL - cm

Define a linear regression model specification.

```
In [8]: # Use `linear_reg()` with engine `lm` and mode `regression`
lm_spec <- linear_reg() %>%
  set_engine("lm") %>%
  set_mode("regression")

#lm_model_weather <- linear_reg(mode = "regression", engine = "lm", penalty =
  NULL, mixture = NULL)
```

Fit a model with the response variable `RENTED_BIKE_COUNT` and predictor variables `TEMPERATURE + HUMIDITY + WIND_SPEED + VISIBILITY + DEW_POINT_TEMPERATURE + SOLAR_RADIATION + RAINFALL + SNOWFALL`

```
In [9]: # Fit the model called `lm_model_weather`
# RENTED_BIKE_COUNT ~ TEMPERATURE + HUMIDITY + WIND_SPEED + VISIBILITY + DEW_P
POINT_TEMPERATURE + SOLAR_RADIATION + RAINFALL + SNOWFALL, with the training d
ata

lm_model_weather <- lm_spec %>% fit(RENTED_BIKE_COUNT ~ TEMPERATURE + HUMIDITY
+ WIND_SPEED + VISIBILITY + DEW_POINT_TEMPERATURE + SOLAR_RADIATION + RAINFALL
+ SNOWFALL, data = train_data)
```

Print the fit summary for the `lm_model_weather` model.


```
In [10]: print(lm_model_weather)
```

parsnip model object

Fit time: 6ms

Call:

```
stats::lm(formula = RENTED_BIKE_COUNT ~ TEMPERATURE + HUMIDITY +  
  WIND_SPEED + VISIBILITY + DEW_POINT_TEMPERATURE + SOLAR_RADIATION +  
  RAINFALL + SNOWFALL, data = data)
```

Coefficients:

(Intercept)	TEMPERATURE	HUMIDITY
156.71	2399.74	-918.38
WIND_SPEED	VISIBILITY	DEW_POINT_TEMPERATURE
404.47	12.56	-316.92
SOLAR_RADIATION	RAINFALL	SNOWFALL
-444.85	-1764.01	317.78

You should see the model details such as formula, residuals, and coefficients.

TASK: Build a linear regression model using all variables

In addition to weather, there could be other factors that may affect bike rental demand, such as the time of a day or if today is a holiday or not.

Next, let's build a linear regression model using all variables (weather + date/time) in this task.

TODO: Build a linear regression model called `lm_model_all` using all variables `RENTED_BIKE_COUNT ~ .`.

```
In [18]: # Fit the model called `lm_model_all`  
# `RENTED_BIKE_COUNT ~ .` means use all other variables except for the response variable  
  
lm_model_all <- lm_spec %>% fit(RENTED_BIKE_COUNT ~ ., data=train_data)  
  
#lm_model_all <- lm(RENTED_BIKE_COUNT ~ ., data=train_data)
```

Print the fit summary for `lm_model_all`.

```
In [20]: summary(lm_model_all$fit)
```

Call:

```
stats::lm(formula = RENTED_BIKE_COUNT ~ ., data = data)
```

Residuals:

	Min	1Q	Median	3Q	Max
	-1401.45	-218.96	-7.31	199.53	1780.67

Coefficients: (3 not defined because of singularities)

	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	212.20	51.04	4.158	3.26e-05	***
TEMPERATURE	782.66	212.13	3.690	0.000227	***
HUMIDITY	-886.73	99.49	-8.913	< 2e-16	***
WIND_SPEED	31.91	40.27	0.792	0.428169	
VISIBILITY	21.87	20.26	1.079	0.280439	
DEW_POINT_TEMPERATURE	598.39	221.37	2.703	0.006888	**
SOLAR_RADIATION	276.88	41.47	6.677	2.64e-11	***
RAINFALL	-2064.64	143.28	-14.410	< 2e-16	***
SNOWFALL	260.97	103.50	2.522	0.011709	*
`0`	-29.30	34.26	-0.855	0.392515	
`1`	-116.85	33.72	-3.465	0.000533	***
`10`	-237.52	32.74	-7.255	4.48e-13	***
`11`	-247.38	33.85	-7.309	3.02e-13	***
`12`	-208.34	34.39	-6.059	1.45e-09	***
`13`	-191.35	35.04	-5.461	4.90e-08	***
`14`	-192.44	34.44	-5.588	2.39e-08	***
`15`	-109.73	34.40	-3.190	0.001429	**
`16`	23.13	34.00	0.680	0.496431	
`17`	305.55	34.15	8.946	< 2e-16	***
`18`	794.80	34.02	23.364	< 2e-16	***
`19`	522.99	34.25	15.268	< 2e-16	***
`2`	-237.21	33.74	-7.030	2.28e-12	***
`20`	432.00	34.13	12.657	< 2e-16	***
`21`	446.58	34.09	13.100	< 2e-16	***
`22`	342.64	33.85	10.123	< 2e-16	***
`23`	103.81	33.85	3.066	0.002175	**
`3`	-319.87	34.23	-9.346	< 2e-16	***
`4`	-386.29	34.00	-11.361	< 2e-16	***
`5`	-362.72	33.48	-10.834	< 2e-16	***
`6`	-204.12	33.59	-6.076	1.30e-09	***
`7`	106.76	33.58	3.179	0.001486	**
`8`	450.98	32.62	13.823	< 2e-16	***
`9`	NA	NA	NA	NA	
AUTUMN	359.00	20.29	17.694	< 2e-16	***
SPRING	191.37	19.36	9.884	< 2e-16	***
SUMMER	198.14	29.19	6.789	1.24e-11	***
WINTER	NA	NA	NA	NA	
HOLIDAY	-124.42	22.95	-5.422	6.11e-08	***
NO_HOLIDAY	NA	NA	NA	NA	

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 377.9 on 6312 degrees of freedom

Multiple R-squared: 0.6589, Adjusted R-squared: 0.657

F-statistic: 348.4 on 35 and 6312 DF, p-value: < 2.2e-16

Now you have built two basic linear regression models with different predictor variables, let's evaluate which model has better performance,

TASK: Model evaluation and identification of important variables

Now that you have built two regression models, `lm_model_weather` and `lm_model_all`, with different predictor variables, you need to compare their performance to see which one is better.

In this project, you will be asked to use very important metrics that are often used in Statistics to determine the performance of a model:

1. R^2 / R-squared
2. Root Mean Squared Error (RMSE)

R-squared

R squared, also known as the coefficient of determination, is a measure to indicate how close the data is to the fitted regression line. The value of R-squared is the percentage of variation of the response variable (y) that is explained by a linear model.

Root Mean Squared Error (RMSE)

$$RMSE = \sqrt{MSE}$$

As you know, the Mean Squared Error measures the average of the squares of errors, where 'error' is the difference between the actual value (y) and the estimated value (\hat{y}). Another metric that is related to MSE is **Root Mean Squared Error (RMSE)** and is simply the square root of MSE.

We first need to test the `lm_model_weather` and `lm_model_all` models against the test dataset `test_data`, and generate `RENTED_BIKE_COUNT` prediction results.

TODO: Make predictions on the testing dataset using both `lm_model_weather` and `lm_model_all` models

```
In [21]: # Use predict() function to generate test results for `lm_model_weather` and `
lm_model_all`
# and generate two test results dataframe with a truth column:

test_results_weather <- lm_model_weather %>% predict(new_data = test_data) %>%
mutate(truth = test_data$RENTED_BIKE_COUNT)

test_results_all <- lm_model_all %>% predict(new_data = test_data) %>% mutate(
truth = test_data$RENTED_BIKE_COUNT)
```

Warning message in predict.lm(object = object\$fit, newdata = new_data, type = "response"):
“prediction from a rank-deficient fit may be misleading”Warning message in predict.lm(object = object\$fit, newdata = new_data, type = "response"):
“prediction from a rank-deficient fit may be misleading”

NOTE: if you happen to see a warning like : prediction from a rank-deficient fit may be misleading , it may be caused by collinearity in the predictor variables. Collinearity means that one predictor variable can be predicted from other predictor variables to some degree. For example, RAINFALL could be predicted by HUMIDITY .

But dont worry, you will address glmnet models (Lasso and Elastic-Net Regularized Generalized Linear Models) instead of regular regression models to solve this issue and futher improve the model performance.

Next, let's calculate and print the R-squared and RMSE for the two test results

TODO: Use rsq() and rmse() functions to calculate R-squared and RMSE metrics for the two test results

```
In [25]: rsq_weather <- rsq(test_results_weather, truth = truth, estimate = .pred)
rsq_weather
```

.metric	.estimator	.estimate
rsq	standard	0.6690204

```
In [26]: rsq_all <- rsq(test_results_all, truth = truth, estimate = .pred)
rsq_all
```

.metric	.estimator	.estimate
rsq	standard	0.6690204

```
In [28]: rmse_weather <- rmse(test_results_weather, truth = truth, estimate = .pred)
rmse_weather
```

.metric	.estimator	.estimate
rmse	standard	364.4235

```
In [30]: rmse_all <- rmse(test_results_all, truth = truth, estimate = .pred)
rmse_all
```

.metric	.estimator	.estimate
rmse	standard	364.4235

From these tables, you should find that the test results from `lm_model_all` are much better. It means that using both weather and datetime variables in the model generates better prediction results.

Since `lm_model_all` has many predictor variables, let's check which predictor variables have larger coefficients. Variables with larger coefficients in the model means they attribute more in the prediction of `RENTED_BIKE_COUNT`. In addition, since all predictor variables are normalized to the same scale, 0 to 1, we thus can compare their coefficients directly.

You could try building another regression model using the non-normalized `seoul_bike_sharing_converted.csv` dataset, and you would find that the coefficients are much different.

First let's print all coefficients:

```
In [31]: lm_model_all$fit$coefficients
```

(Intercept)	212.199462743954
TEMPERATURE	782.657870088692
HUMIDITY	-886.73027784441
WIND_SPEED	31.9129927279246
VISIBILITY	21.8716449661126
DEW_POINT_TEMPE...	598.387261041619
SOLAR_RADIATION	276.881652368234
RAINFALL	-2064.63754308304
SNOWFALL	260.973302878596
`0`	-29.2983361016137
`1`	-116.846747038285
`10`	-237.51820120511
`11`	-247.38341321939
`12`	-208.341330868407
`13`	-191.354228014138
`14`	-192.441682312028
`15`	-109.733640240686
`16`	23.128163602844
`17`	305.547618249471
`18`	794.803151569959
`19`	522.988214834461
`2`	-237.211663999792
`20`	431.99523757542
`21`	446.580533576852
`22`	342.64091922177
`23`	103.808410131075
`3`	-319.871517124491
`4`	-386.292807947984
`5`	-362.720001875814
`6`	-204.118732579813
`7`	106.757133838754
`8`	450.977534093274
`9`	<NA>
AUTUMN	358.999443905867
SPRING	191.365267451116
SUMMER	198.142010192369
WINTER	<NA>
HOLIDAY	-124.423741444013
NO_HOLIDAY	<NA>

hmm, it's not very clear to compare the coefficients from a long and unsorted list. Next, you need to sort and visualize them using a bar chart

TODO: Sort the coefficient list in descending order and visualize the result using `ggplot` and `geom_bar`

```
In [41]: sorted <- sort(abs(lm_model_all$fit$coefficients), decreasing = TRUE)
sorted
```


RAINFALL	2064.63754308304
HUMIDITY	886.73027784441
`18`	794.803151569959
TEMPERATURE	782.657870088692
DEW_POINT_TEMPE...	598.387261041619
`19`	522.988214834461
`8`	450.977534093274
`21`	446.580533576852
`20`	431.99523757542
`4`	386.292807947984
`5`	362.720001875814
AUTUMN	358.999443905867
`22`	342.64091922177
`3`	319.871517124491
`17`	305.547618249471
SOLAR_RADIATION	276.881652368234
SNOWFALL	260.973302878596
`11`	247.38341321939
`10`	237.51820120511
`2`	237.211663999792
(Intercept)	212.199462743954
`12`	208.341330868407
`6`	204.118732579813
SUMMER	198.142010192369
`14`	192.441682312028
SPRING	191.365267451116
`13`	191.354228014138
HOLIDAY	124.423741444013
`1`	116.846747038285
`15`	109.733640240686
`7`	106.757133838754
`23`	103.808410131075
WIND_SPEED	31.9129927279246
`0`	29.2983361016137
`16`	23.128163602844
VISIBILITY	21.8716449661126

```
In [111]: sorted_df <- as.data.frame(sorted)
sorted_df$coefficient <- as.numeric(sorted_df$sorted)
sorted_df$feature <- row.names(sorted_df)
sorted_df
```



	sorted	coefficient	feature
RAINFALL	2064.63754	2064.63754	RAINFALL
HUMIDITY	886.73028	886.73028	HUMIDITY
`18`	794.80315	794.80315	`18`
TEMPERATURE	782.65787	782.65787	TEMPERATURE
DEW_POINT_TEMPERATURE	598.38726	598.38726	DEW_POINT_TEMPERATURE
`19`	522.98821	522.98821	`19`
`8`	450.97753	450.97753	`8`
`21`	446.58053	446.58053	`21`
`20`	431.99524	431.99524	`20`
`4`	386.29281	386.29281	`4`
`5`	362.72000	362.72000	`5`
AUTUMN	358.99944	358.99944	AUTUMN
`22`	342.64092	342.64092	`22`
`3`	319.87152	319.87152	`3`
`17`	305.54762	305.54762	`17`
SOLAR_RADIATION	276.88165	276.88165	SOLAR_RADIATION
SNOWFALL	260.97330	260.97330	SNOWFALL
`11`	247.38341	247.38341	`11`
`10`	237.51820	237.51820	`10`
`2`	237.21166	237.21166	`2`
(Intercept)	212.19946	212.19946	(Intercept)
`12`	208.34133	208.34133	`12`
`6`	204.11873	204.11873	`6`
SUMMER	198.14201	198.14201	SUMMER
`14`	192.44168	192.44168	`14`
SPRING	191.36527	191.36527	SPRING
`13`	191.35423	191.35423	`13`
HOLIDAY	124.42374	124.42374	HOLIDAY
`1`	116.84675	116.84675	`1`
`15`	109.73364	109.73364	`15`
`7`	106.75713	106.75713	`7`
`23`	103.80841	103.80841	`23`
WIND_SPEED	31.91299	31.91299	WIND_SPEED
`0`	29.29834	29.29834	`0`
`16`	23.12816	23.12816	`16`
VISIBILITY	21.87164	21.87164	VISIBILITY

```
In [121]: sorted2 <- sorted_df                                     # Replica
           te original data
sorted2$feature <- factor(sorted2$feature,
                           # Factor levels in increasing order
                           levels = sorted2$feature[order(sorted2$coefficient)])

ggplot(data=sorted2,aes(feature,coefficient))+
  geom_bar(stat = "identity")+
  coord_flip()
```



You should see a sorted coefficient bar chart like the following example:



(https://cognitiveclass.ai/?utm_medium=Exinfluencer&utm_source=Exinfluencer&utm_content=000026UJ&utm_term=10006555&utm_id=NA&SkillsNetwork-Channel-SkillsNetworkCoursesIBMDveloperSkillsNetworkRP0321ENSkillsNetwork25371262-2021-01-01)

Mark down these 'top-ranked variables by coefficient', which will be used for model refinements in the next labs.

Next Steps

Great! Now you have built a baseline linear regression model to predict hourly bike rent count, with reasonably good performance. In the next lab, you will be refining the baseline model to improve its performance.

Authors

[Yan Luo \(https://www.linkedin.com/in/yan-luo-96288783/?utm_medium=Exinfluencer&utm_source=Exinfluencer&utm_content=000026UJ&utm_term=10006555&utm_id=NA&utm_campaign=SkillsNetwork-Channel-SkillsNetworkCoursesIBMDeveloperSkillsNetworkRP0321ENSkillsNetwork25371262-2021-01-01\)](https://www.linkedin.com/in/yan-luo-96288783/?utm_medium=Exinfluencer&utm_source=Exinfluencer&utm_content=000026UJ&utm_term=10006555&utm_id=NA&utm_campaign=SkillsNetwork-Channel-SkillsNetworkCoursesIBMDeveloperSkillsNetworkRP0321ENSkillsNetwork25371262-2021-01-01)

Other Contributors

Jeff Grossman

Change Log

Date (YYYY-MM-DD)	Version	Changed By	Change Description
2021-04-08	1.0	Yan	Initial version created

© IBM Corporation 2021. All rights reserved.